

Reinforcement Learning

Michael Plüss
Elias Arnold

2016-051-062 (M. Plüss)
2014-930-770 (E. Arnold)
m.pluess@unibas.ch
elias.arnold@stud.unibas.ch

1 Introduction

Reinforcement learning attempts to solve the problem of the agent's behaviour in the absence of ground truth data. Such an agent improves its strategy over time and does not need an initial training step. The basic idea is that we feed our agent with a reward based on the agents decisions [Figure 1]. Initially, the agent applies actions just with trial-and-error, and then, from these observed rewards, the agent tries to learn an optimal strategy for the environment. In other learning strategies like supervised learning, the agent updates its behaviour based on the ground truth labels from a supervisor. In reinforcement learning, a supervisor is not available anymore. It just replaces these ground truth labels with the rewards from the environment. Remember, that in unsupervised learning neither ground truth data nor rewards were given.

Reinforcement learning is a term from psychology and reflects the learning behaviour of nature. For example, animals learn to behave in such a way that hunger is minimized while maximizing hunting success. There are many applications of reinforcement learning in the realm of computer science like image recognition, robotics and computer game AI's.

In this summary of chapter 21 from [1], we will first look at the different approaches to designing an agent and then discuss specific methods to implement reinforcement learning. Next we will generalize the insights from the specific methods. Finally, we will take a look at dynamically modifying the strategy, under which the agent operates.




Fig. 1. The basic structure of a reinforcement learning agent. The agent finds itself in state S_t and applies action A_t . The resulting state S_{t+1} and reward R_{t+1} are fed back to the agent [2].

2 Agent Design

Depending on the environment we have given, the agent might or might not know the transition model or the reward function of the environment. In reinforcement learning, we assume the case where the agent has no prior knowledge about any of these things. In order to be able to determine suitable transitions, the agent sometimes has to use certain information about the world. The following three agents look at different kinds of information.

Throughout this summary, we will refer to such a strategy of the agent as a policy. A policy $\pi(s) \rightarrow a$ is a mapping between a state in the environment s and an applicable action a . So it tells the agent which action to execute in a certain state.

2.1 Utility-based Agent

A utility-based agent takes a state as input and calculates the utility solely from the state. It tries to take the action which moves it to the state with the highest expected utility. Thus, it needs to know which actions are available and to which state these actions will lead. In other words, it needs to learn a model of the environment.

2.2 Q-learning Agent

A Q-learning agent learns an action-utility function (or Q-function). This means it learns the expected utility of taking a given action in a given state. Such an agent can compare the usefulness of actions (Q-value) without having to know where these actions take it and thus it does not need to learn a model of the environment. Therefore, we call this agent “model-free”. A drawback of such an agent is, that it cannot look ahead, because it does not know the successor state where its applied actions lead it to.

2.3 Reflex Agent

A reflex agent decides what to do based solely on the current state. It reacts to its environment directly. As such, it does not require a model of the world or the utilities of states or actions. The downside is that such an agent is normally restricted to simple problems. An example would be the AI in a pong game, where it just moves into the direction of the ball. This can be seen as a policy and the agent is therefore “policy-based”.

3 Types of Learning

The goal of the agent is to indicate for each state what the expected reward of a path that leads through the state might be. This is called a utility of a state $U(s)$. The methods to learn these utilities can be divided into passive and active learning.

3.1 Passive Learning

In passive learning, the policy is given and fixed. This means that it is predetermined which action the agent should take in each state. In such a situation, we might want to learn the utilities of states or state-action pairs. The goal of passive learning is to label each state with its utility, such that the Bellman equation holds in each state. In contrast to policy evaluation, we assume the action outcome probabilities $P(s'|s, a)$ to be unknown, so we can not just apply the Bellman equation to calculate a states utility.

$$U^\pi(s) = R(s) + \sum_{s'} P(s'|s, \pi(s)) U^\pi(s')$$

The Bellman equation, where s is the current state, a the applied action, and s' the state in which the agent ends up in. Passive learning just tries to make this equilibrium hold in each state.

3.2 Active Learning

In active learning, the agent must additionally find out what to do in each state because the policy is not given anymore. In passive learning, a policy π assigns an action a to each state s where the mapping was independent of the states that were already visited. In the active case, the agent must reassess after each transition which action to choose next. The path that the agent has already completed plays an important role in these decisions. One of the main challenges with active learning, is to factor in a certain amount of exploration, where the agent tries to take actions, which were not taken before. The hope is, that by doing so, the agent reaches previously unseen states and possibly discovers a shortcut or a better solution. This topic is discussed in detail in chapter 5.

In active learning, the utilities of the individual states must converge to the Bellman equation too. However, the equation must be adjusted to take into account all executable actions, not just the one which is decided by the policy like in passive learning.

$$U(s) = R(s) + \gamma \max_a \sum_{s'} P(s'|s, a) U(s')$$

4 Algorithms

So far, we looked at different strategies for the agent to gather the information needed to calculate the current state's utility. We have seen, that these methods can be split up in active and passive learning methods. In this section, we will look at different algorithms to actually calculate the utility values. For each algorithm we will indicate, whether it can be used in active and passive learning.

4.1 Direct Utility Estimation

The idea of direct utility estimation (DUE) is to let the agent run several trials and to gradually modify the utility of each visited state after each trial. We use a fixed policy for running the trials and thus, DUE is an algorithm for passive learning. In the longer run, the utility estimate converges to the true values (so that the Bellman equation holds in each state).

After each trial, we calculate the utility of each visited state s based on the rewards received by the following states S_t with the following formula:

$$U^\pi(s) \leftarrow E[\sum_{t=0}^{\infty} \gamma^t R(S_t)] \text{ where } \gamma \text{ denotes a discount factor and } S_0 = s.$$

We do not directly assign the calculated values $U(s)$ to s however. Rather, we keep a running average, that we modify after each trial. This algorithm is very simple, but has a few drawbacks: Firstly, it is not guaranteed to visit every state in the environment (and thus it does not calculate the state's utilities). Secondly, this algorithm ignores dependencies between neighbouring states, which are included in the bellman equation via $P(s'|s, a)$. Consequently, this algorithm converges rather slowly to the true utility values, when compared with other algorithms. However, the absence of these action-outcome probabilities $P(s'|s, a)$ can also be seen as an advantage, because these values are rather hard to estimate from trials.

4.2 Adaptive Dynamic Programming

Another algorithm for the passive learning scenario is Adaptive Dynamic Programming (ADP). ADP works similar to direct utility estimation in the way that it also lets the agent run through many trials in order to estimate the utility of the states. The difference to the first approach is, that ADP estimates the transition model $P(s'|s, a)$ from example trials.

It achieves this by keeping track of the number of times a state s has been visited and which action a we execute from there in the array N_{sa} and in which state s' we land from there in the array $N_{s'|sa}$. From the ratio of these values we can calculate the probability of reaching state s' from s with the action a as $P(s'|s, a) = \frac{N_{s'|sa}}{N_{sa}}$. This value will approximate the true transition model, given enough trials. Finally, to calculate the utility of each state, we plug our estimated transition model into the Bellman equation for passive learning (with a fixed policy π):

$$U^\pi(s) \leftarrow R(s) + \gamma \sum_{s'} [P(s'|s, \pi(s)) U^\pi(s')]$$

In the formula γ again denotes a discount factor.

We can modify this algorithm to use ADP in an active learning scenario. In that case, we just plug the estimated transition model in the Bellman equation for the active learning case. The resulting formula looks as follows:

$$U(s) \leftarrow R(s) + \gamma \max_a \sum_{s'} [P(s'|s, a) * U(s')]$$

Effectively, here we are looking at each possible action a that we can take in a state s and choose the action, which promises the best outcome. This is required, since we do not have a fixed policy anymore. Such a one-step look-ahead approach provides an optimal policy.

4.3 Temporal Difference Learning

Temporal-difference learning (TD) is another algorithm to estimate the utility of states in a passive learning scenario.

For TD we again use a large number of trials. The idea is, as before in the direct utility estimation, to calculate the utility of a state based on its successor. This time we use the difference of utilities between successive states together with a learning parameter α , which becomes smaller as a state is visited more often. The new formula to repeatedly modify the utility values of the states looks like this:

$$U^\pi(s) \leftarrow U^\pi(s) + \alpha(N_s) * (R(s) + \gamma U^\pi(s') - U^\pi(s))$$

At first glance, this formula looks confusing. If we ignore the learning rate α , the equation can be rewritten as $U^\pi(s) = R(s) + U^\pi(s')$. Now the closeness to Adaptive Dynamic Programming can be seen. So in TD we just add $R(s)$ with the utility of the successor $U(s')$, while ADP calculates the expected value of the utility of s' . γ again denotes a discount factor. This algorithm behaves in a similar fashion as the ADP, but requires significantly less computational power.

4.4 Learning State-Action-Utilities (Q-Learning)

Another method that is very similar to Temporal Difference Learning, but works in an active learning scenario, is Q-learning. This method learns the utility of state-action pairs ($Q(s, a)$) instead of just the state utilities ($U(s)$), which would not provide any benefits in the passive learning case, since there we have the policy given.

The connection between the passive Temporal Difference Learning and the active Q-Learning is given by $U(s) = \max_a Q(s, a)$, where $Q(s, a)$ is the utility of action a given state s . This allows us to rewrite the Temporal Difference update rule as:

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \gamma * \max_{a'} Q(s', a') - Q(s, a))$$

α is the learning rate and γ is the discount factor. Note, that when we are in a state s and apply action a , we consider all possibly following states s' and actions a' that can be applied next. In other terms, Q-Learning rates a state-action pair only by the best following pair.

4.5 The SARSA Algorithm

For the active learning case, there is another alternative to Q-Learning, which also learns Q-values. This method is called SARSA, which is an abbreviation for State-Action-Reward-State-Action (s, a, r, s', a') . This algorithm does not consider all pairs (s', a') (like Q-Learning), but waits until an action a' in state s' has actually been executed:

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \gamma Q(s', a') - Q(s, a))$$

Because SARSA needs to know in each state s what the following state s' will be, this algorithm is called “on-policy”, while other Q-learning algorithms which do not need a policy are called “off-policy”. SARSA converges to the optimal policy, but at a much slower speed than ADP. This algorithm is more realistic than Q-Learning, because it updates a state’s Q-value only if s' and action a' are actually visited/executed. (Same difference as between ADP and TD in the passive case)

5 Balancing Exploration vs. Exploitation

With the introduction of active learning, wherein we do not follow a fixed policy but decide where to go before each step of the agent, a new issue emerges: Upon finding a solution, a greedy agent will exploit this solution and steadily fortify the found path. The agent will thus miss the chance to take another route and possibly find a shortcut. The cure to this problem is the introduction of a kind of curiosity via exploration.

The basic idea is, that an agent first tries out many different actions, preferring those which it did not take previously. After a while, when the agent has explored the state-space to a reasonable extent, it will shift its behaviour from exploration towards exploitation. Such an approach is called greedy in the limit of infinite exploration (GLIE).

As you may suspect, the tricky part is to find the right time when to switch from exploring to exploiting.

One way to implement such a GLIE scheme is to introduce a probability to take a random action instead of the greedy one. This function must decay over time to fit the GLIE constraints, for example $\frac{1}{t}$.

Alternatively, you can assign a reward function f to a state instead of a constant reward. This function gives a state a reward (makes the state more attractive) depending on the number of times this state was already visited. A possible implementation would be:

$$f(u, n) = \begin{cases} R^+ & \text{if } n < N_e \\ u & \text{otherwise} \end{cases}$$

Where R^+ is an optimistic estimate of the best possible reward that any state can have, n indicates how often a state has been visited, N_e is some constant

threshold, and u is the classical utility of a state according to the Bellman equation.

This formula makes unvisited states artificially appear attractive to the pathfinding algorithm (with R^+). But if after N_e attempts via this (initially unexplored) state no better trial is found, the state is assigned with his original reward.

A possible application of this function is for example to modify the update rule of the ADP algorithm, which then becomes:

$$U^+(s) \leftarrow R(s) + \gamma \max_a f\left(\sum_{s'} [P(s'|s, a) * U^+(s')], N(s, a)\right)$$

By using U^+ we signify, that we use a version of the ADP algorithm, wherein we use the modified rewards of states.

6 Generalization in Reinforcement Learning

In the previous chapters we discussed methods to calculate state(-action)-utilities for every state the agent can end up in. We had to consider all possible successor states/actions to come up with the true utility estimate for a state. While ADP converged to the optimal policy relatively quickly, Temporal Difference Learning, and SARSA take much longer. Since we do not have a model in these cases, we have to consider a lot sample trials for the optimal utility estimate.

These approaches are reasonable for small state-spaces. On the contrary, if the state-space were large, it would no longer be justifiable to calculate the utility estimate for each state in this way. Therefore, it is often attempted to approximate the true utility of a state with a function that can be evaluated in one sweep. Such a function approximation has the following form:

$$\hat{U}_\theta(s) \leftarrow \theta_1 f_1(s) + \theta_2 f_2(s) + \dots + \theta_n f_n(s)$$

The f s are called features and the θ s are called the parameters of the evaluation function $\hat{U}_\theta(s)$. The main benefit of function approximation is the ability to generalize from a few states where the true utility estimate is known to a larger state-space where we cannot calculate the utility of every single state. For example, the state-space of the game backgammon contains 10^{20} states. Tesauro [3] showed in 1992 that a computer program could play backgammon as well as a human, with a function approximation that could only model 10^{12} different states. The f s are fixed functions of the state variables, that are weighted by their respective θ s. These θ s need to be adjusted after each trial (idea of reinforcement learning). For example, we can use the Widrow-Hoff rule (also called “Delta rule”) to update the parameter θ_i :

$$\theta_i \leftarrow \theta_i - \alpha \frac{\partial E_j(s)}{\partial \theta_i} \text{ where } E_j(s) = \frac{(\hat{U}_\theta(s) - u_j(s))^2}{2} \text{ (MSE)}$$

In the above equation, $u_j(s)$ denotes the observed utility of a state s in the j -th trial and $\hat{U}_\theta(s)$ denotes the predicted utility, according to formula ???. The f s are usually functions of the form $f_i(x, y)$, where x, y are features of the state, that are weighted by θ_i . A common choice is a function that returns the agents closeness to the state with the highest reward. Let us consider a grid world of the size 100^2 and the state with reward $+1$ has the coordinates $(50, 50)$. Then, $\hat{U}_\theta(s)$ may look like:

$$\hat{U}_\theta(s) = \theta_1 f_1(x, y) = \theta_1 \sqrt{(x - 50)^2 + (y - 50)^2}$$

Here, the utility of a state depends solely on the distance to the goal state, weighted with θ_1 .

7 Policy Search

The general idea of policy search is to change the policy for as long as there is improvement and then stop. Normally, we want to work with policies π that are parametrized. We apply the changes to the parameters θ . Usually, such policies have a lot fewer parameters than there are states in the state-space of the problem. This ensures a certain degree of generalization and also makes it much easier to mathematically work with the policies. As an example, we could define the obvious policy which always takes the action a from the state s that predicts the biggest reward. This would result in a greedy agent, that never explores unseen states.

7.1 Problems

A problem that we have to be aware about when using policy search, is that the policy may be a discontinuous function. The problem with that is that small changes in θ may lead to very different outcomes. To illustrate this, let us picture an example where a small change of θ leads an agent to take the left path, instead of the right one. Because the agent took another path at the intersection, it now lands in another goal state with a considerably higher reward. This jumpy behaviour of the policy function can make it very difficult or even impossible to take its derivative, which we need to perform optimization techniques such as hill-climbing. Luckily, we can tackle this problem by using a stochastic policy representation $\pi_\theta(s, a)$, which gives us a probability of selecting action a in state s . An often-used approach to achieve this is the softmax function:

$$\pi_\theta(s, a) = e^{Q_\theta(s, a)} / \sum_{a'} e^{Q_\theta(s, a')}$$

7.2 Policy Improvement

How do we now actually improve the policy? To evaluate how good our policy actually is, we need a function $\rho(\theta)$, called the policy value, which represents the

expected reward for executing the policy π_θ . If we are able to find an expression for $\rho(\theta)$ in closed form, this becomes a standard optimization problem and we can simply follow the policy gradient $\nabla_\theta \rho(\theta)$, provided $\rho(\theta)$ is differentiable, but oftentimes this is not possible.

An alternative is to execute the policy π_θ and take the accumulated reward of the trial as the value for $\rho(\theta)$. In this situation, we can observe the change of $\rho(\theta)$ for small changes of θ and apply the change which results in the biggest improvement of $\rho(\theta)$.

When we have a situation where the environment or the policy is stochastic, this becomes a bit more complicated. Executing two trials with the same policy may lead to very different outcomes, because the path that the agent takes is not deterministic anymore. The solution is to run several trials and take the average of the accumulated rewards as $\rho(\theta)$. The drawback of this approach is, that it may be very expensive, time-consuming or even dangerous to run a trial, so this is not always an option. Luckily it is possible to obtain an unbiased estimate of the gradient $\nabla_\theta \rho(\theta)$ at θ . We will assume a nonsequential environment in which the reward $R(a)$ is collected directly after taking action a in the start state s_0 . Here, $R(a)$ stands for the sum of subsequent rewards in the trial(s) where we chose action a . This algorithm is called “REINFORCE” (Williams, 1992)[5] and is usually much more efficient than standard hill-climbing.

7.3 Conclusion

In this chapter, we have learnt how to use reinforcements (rewards) to learn utilities or policies. For that we have looked at different ways to design agents (focusing on the utilities of states vs. state-action pairs) and different algorithms to compute the utilities. These algorithms can be used in different settings: In a passive setting where the policy is given and fixed, or in an active scenario, where the policy is not given and has to be found.

Furthermore we have looked at the issue of exploration vs. exploitation, where the difficulty lies in balancing the curious exploration of unknown states vs. greedily moving towards the most attractive state.

We also discussed the generalization of utility and/or policy calculation, such that we can use reinforcement learning with big state-spaces, where the previously discussed methods would not be feasible. Learning parameters of base functions also enables the abstraction of a state’s utility based on its features. This has the advantage, that we can already estimate the utilities of states that we see for the first time.

In practice, reinforcement learning has proven to be a viable learning strategy. Reinforcement learning has applications in robotics, games and chemistry, making it a very flexible concept with veritable benefits.

“There is no reasoning, no process of inference or comparison; there is no thinking about things, no putting two and two together; there are no ideas-the

animal does not think of the box or of the food or of the act he is to perform.”
- Edward Thorndike [4], the psychologist who proposed Law of effect.

References

1. Russell, S., Norvig, P.: Artificial Intelligence: A Modern Approach. Prentice Hall Press, Upper Saddle River, NJ, USA, 3rd edn. (2009)
2. Sutton, R.S., Barto, A.G.: Reinforcement learning: An introduction. *IEEE Trans. Neural Networks* 9(5), 1054–1054 (1998), <https://doi.org/10.1109/TNN.1998.712192>
3. Tesauro, G.: Practical issues in temporal difference learning. *Mach. Learn.* 8(3-4), 257–277 (May 1992), <https://doi.org/10.1007/BF00992697>
4. Thorndike, E.: Animal Intelligence: Experimental Studies. Hafner Publishing Company (1965), <https://books.google.de/books?id=0uPaAAAAMAAJ>
5. Williams, R.J.: Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine Learning* 8(3), 229–256 (May 1992), <https://doi.org/10.1007/BF00992696>